

Question 1 — What is Anomaly Detection?

Paragraph:

Anomaly detection is the process of identifying observations, events, or data points that deviate significantly from the expected pattern in a dataset. These anomalies may indicate errors, fraud, rare events, or novel phenomena and are important because they often signal issues or opportunities requiring attention.

Points:

- **Point anomalies:** single data points that are abnormal (e.g., a taxi fare of \$10,000 in a city dataset).
 - **Contextual anomalies:** points anomalous in context (e.g., temperature 30°C in winter is anomalous in context but not globally).
 - **Collective anomalies:** a group of related observations abnormal together (e.g., a sudden drop in energy usage across many sensors indicating power outage).
-

Question 2 — Compare Isolation Forest, DBSCAN, and Local Outlier Factor

Paragraph:

These three methods take different approaches: Isolation Forest isolates anomalies by random partitioning, DBSCAN clusters density and flags points in low-density regions as outliers, and Local Outlier Factor (LOF) compares local density of a point to its neighbors to assess outlierness. Each is suitable in different contexts depending on dimensionality, density structure, and whether labeled data or cluster structure is expected.

Points (short):

- **Isolation Forest**
 - **Approach:** Tree-based random partitioning; anomalies require fewer splits.
 - **Strengths:** Fast, scales to high dimensions, unsupervised.
 - **Use cases:** High-dimensional numeric data, fraud detection, telemetry.
- **DBSCAN**
 - **Approach:** Density-based clustering; points not in dense clusters → noise/outliers.
 - **Strengths:** Finds arbitrary-shaped clusters, identifies noise.
 - **Limitations:** Sensitive to `eps` and `min_samples`; struggles with varying density and very high dimensions.

- **Use cases:** Spatial data, geolocation clustering, cases with clear density separation.
 - **Local Outlier Factor (LOF)**
 - **Approach:** Compares local density of a point to densities of its neighbors (k-NN).
 - **Strengths:** Detects local anomalies where global methods fail.
 - **Limitations:** Sensitive to k ; $O(n^2)$ naive for large n (but scikit-learn is optimized).
 - **Use cases:** Datasets with varying density, localized anomaly detection.
-

Question 3 — Key components of a Time Series

Paragraph:

A time series is typically decomposed into components that explain its behavior: trend (long-term increase/decrease), seasonality (repeating patterns), and residual/noise (unexplained variation). Understanding these helps modeling, forecasting, and anomaly detection.

Points (one example each):

- **Trend:** long-term direction (e.g., steadily increasing monthly customers).
 - **Seasonality:** repeating patterns at fixed intervals (e.g., higher electricity demand every summer).
 - **Cyclic (optional):** non-fixed periodic fluctuations (e.g., business cycles).
 - **Residual / Noise:** random variations after removing trend & seasonality (e.g., measurement error).
-

Question 4 — Define Stationary & how to test/transform

Paragraph:

A stationary time series has a constant mean, variance, and autocovariance over time (statistical properties do not change with time). Many time series models require stationarity because model parameters assume stable relationships.

Points (tests & transforms):

- **Tests for stationarity:**
 - Augmented Dickey–Fuller (ADF) test (null: unit root / non-stationary).

- KPSS test (null: stationary).
 - Visual inspection of rolling mean/variance and ACF plots.
 - **Transforms to induce stationarity:**
 - Differencing (first or seasonal differences).
 - Log or Box–Cox transform (stabilize variance).
 - Seasonal differencing (difference lagged by seasonal period).
 - **Workflow:** test → transform (e.g., **log** then **diff**) → retest until stationary.
-

Question 5 — AR, MA, ARIMA, SARIMA, SARIMAX (structure & use)

Paragraph:

Autoregressive (AR) models use past values to predict the present; Moving Average (MA) models use past errors. ARIMA combines AR and MA with differencing to handle non-stationary series. SARIMA adds seasonal ARIMA terms for seasonal patterns. SARIMAX extends SARIMA by allowing exogenous regressors (external variables).

Points (short):

- **AR(p):** depends on **p** lagged values. Use when autocorrelation present.
 - **MA(q):** depends on **q** lagged forecast errors. Use when moving-average structure in residuals.
 - **ARIMA(p,d,q):** AR + differencing (**d**) + MA. Use for non-seasonal, non-stationary series.
 - **SARIMA(p,d,q)(P,D,Q,s):** ARIMA with seasonal (**P, D, Q**) terms and period **s**. Use for monthly/quarterly seasonal data.
 - **SARIMAX:** SARIMA + exogenous variables (e.g., weather, holidays). Use when external regressors affect the series.
-

Question 6 — Load AirPassengers, plot series, and decompose

Paragraph:

Below is a compact Python snippet to load the classic Airline Passengers series (AirPassengers), plot the original monthly series, and decompose into trend, seasonal, and residual components using **seasonal_decompose**. If the dataset is not available in your environment, the snippet falls back to reading a local CSV named **AirPassengers.csv**.

```

# Q6: Load AirPassengers, plot and decompose
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.seasonal import seasonal_decompose

# Try several ways to load AirPassengers (robust to environment)
try:
    # preferred if statsmodels provides it
    from statsmodels.datasets import airpassengers
    data = airpassengers.load_pandas().data
    # data has 'value' column or similar - adapt if needed
    if 'value' in data.columns:
        ts = data['value']
    else:
        ts = data.iloc[:, 0]
    ts.index = pd.date_range(start='1949-01', periods=len(ts), freq='MS')
except Exception:
    try:
        # fallback: read local CSV expected in working directory
        df = pd.read_csv('AirPassengers.csv')
        # assume columns: Month, Passengers
        df['Month'] = pd.to_datetime(df['Month'])
        df = df.set_index('Month')
        ts = df.iloc[:,0].asfreq('MS')
    except Exception:
        # final fallback: generate a synthetic seasonal series (for demo)
        idx = pd.date_range('2000-01', periods=120, freq='MS')
        import numpy as np
        ts = pd.Series(100 + 10 * np.sin(2 * np.pi * idx.month / 12) +
            np.linspace(0, 20, len(idx)),
            index=idx)

# plot original series
plt.figure(figsize=(10,4))
plt.plot(ts, linewidth=1.5)
plt.title('AirPassengers (original series)')
plt.xlabel('Date')
plt.ylabel('Passengers')
plt.tight_layout()
plt.show()

# decompose (additive or multiplicative depending on nature; AirPassengers
is multiplicative)
decomp = seasonal_decompose(ts, model='multiplicative', period=12)
# plot decomposition
decomp.plot()

```

```
plt.tight_layout()
plt.show()
```

Notes: adapt column names if your CSV differs. The `period=12` is for monthly seasonality.

Question 7 — Isolation Forest on NYC Taxi Fare (detect & plot anomalies)

Paragraph:

This snippet attempts to load a `nyc_taxi.csv` file with numeric features (e.g., `fare_amount`, `trip_distance`) and runs `IsolationForest`. If no file present, it creates a synthetic 2D dataset. The code flags anomalies and visualizes them on a 2D scatter plot.

```
# Q7: Isolation Forest on NYC Taxi Fare (2D scatter)
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.ensemble import IsolationForest

# load dataset or create synthetic demo
try:
    df = pd.read_csv('nyc_taxi.csv') # user should have this file with
    columns like fare_amount, trip_distance
    X = df[['fare_amount', 'trip_distance']].dropna()
except Exception:
    # synthetic demo (most likely fine for visualization)
    rng = np.random.RandomState(42)
    X_inliers = rng.normal(loc=[10, 3], scale=[3,1], size=(300,2))
    X_outliers = rng.uniform(low=[30,0], high=[80,8], size=(12,2))
    X = pd.DataFrame(np.vstack([X_inliers, X_outliers]),
        columns=['fare_amount', 'trip_distance'])

# fit Isolation Forest
clf = IsolationForest(contamination=0.03, random_state=42)
clf.fit(X)
scores = clf.decision_function(X)
anomaly_flags = clf.predict(X) # -1 anomaly, 1 normal

# plot
plt.figure(figsize=(8,5))
normal = X[anomaly_flags == 1]
anom = X[anomaly_flags == -1]
plt.scatter(normal.iloc[:,0], normal.iloc[:,1], s=25, label='normal',
    alpha=0.6)
```

```
plt.scatter(anom.iloc[:,0], anom.iloc[:,1], s=60, label='anomaly',
edgecolor='k')
plt.xlabel('fare_amount')
plt.ylabel('trip_distance')
plt.title('IsolationForest - anomalies in fare vs distance')
plt.legend()
plt.show()
```

Points:

- Choose `contamination` roughly if you have no ground truth.
- Scale features if they differ in magnitude (e.g., `StandardScaler`).
- For many features, visualize via PCA or t-SNE into 2D.

Question 8 — Train SARIMA on airline passengers and forecast 12 months

Paragraph:

Use `statsmodels` SARIMAX to model the seasonal airline passenger series and forecast the next 12 months. Below code fits a simple SARIMA with seasonal order (1,1,1,12); in practice use AIC/BIC or `pmdarima`'s `auto_arima` to select orders.

```
# Q8: SARIMA train & forecast next 12 months
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.statespace.sarimax import SARIMAX

# Load as in Q6 (reuse ts)
# If running in same notebook, ts is already defined. Otherwise, re-load
the AirPassengers series.
try:
    ts # if defined above
except NameError:
    # fallback loader similar to Q6
    from statsmodels.datasets import airpassengers
    ts = airpassengers.load_pandas().data.iloc[:,0]
    ts.index = pd.date_range(start='1949-01', periods=len(ts), freq='MS')

# ensure float and no missing
ts = ts.astype(float).asfreq('MS')

# fit SARIMAX -- simple baseline orders
```

```

model = SARIMAX(ts, order=(1,1,1), seasonal_order=(1,1,1,12),
enforce_stationarity=False, enforce_invertibility=False)
res = model.fit(dispatch=False)

# forecast next 12 months
n_forecast = 12
pred = res.get_forecast(steps=n_forecast)
pred_ci = pred.conf_int()

# plotting
plt.figure(figsize=(10,5))
plt.plot(ts, label='observed')
pred_index = pd.date_range(ts.index[-1] + pd.offsets.MonthBegin(1),
periods=n_forecast, freq='MS')
plt.plot(pred_index, pred.predicted_mean, label='forecast', marker='o')
plt.fill_between(pred_index, pred_ci.iloc[:,0], pred_ci.iloc[:,1],
alpha=0.2)
plt.title('SARIMA Forecast (next 12 months)')
plt.legend()
plt.show()

# print summary
print(res.summary())

```

Notes:

- Use `auto_arima` (pmdarima) for order selection in real work.
- Validate with rolling-origin cross-validation.

Question 9 — Local Outlier Factor (LOF) and visualize

Paragraph:

LOF compares a point's local density to its neighbors; values much greater than 1 indicate outliers. The snippet loads a dataset or creates synthetic 2D data, applies LOF, and plots anomalies using matplotlib.

```

# Q9: LOF anomaly detection and plot
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.neighbors import LocalOutlierFactor

# load or create 2D data

```

```

try:
    df = pd.read_csv('some_numeric_dataset.csv') # replace with real file
    if available
        X = df[['feature1', 'feature2']].dropna()
except Exception:
    # synthetic mixture
    rng = np.random.RandomState(0)
    X_in = 0.3 * rng.randn(200, 2)
    X_out = rng.uniform(low=-4, high=4, size=(15,2))
    X = pd.DataFrame(np.vstack([X_in, X_out]),
columns=['feature1', 'feature2'])

# fit LOF (novelty=False: fits and predicts in same data)
lof = LocalOutlierFactor(n_neighbors=20, contamination=0.05)
y_pred = lof.fit_predict(X) # -1 = outlier, 1 = inlier
scores = -lof.negative_outlier_factor_ # higher = more outlying

# plot
plt.figure(figsize=(8,5))
inliers = X[y_pred == 1]
outliers = X[y_pred == -1]
plt.scatter(inliers.iloc[:,0], inliers.iloc[:,1], s=25, label='inlier',
alpha=0.6)
plt.scatter(outliers.iloc[:,0], outliers.iloc[:,1], s=80, label='outlier',
edgecolor='k')
plt.legend()
plt.title('Local Outlier Factor - detected anomalies')
plt.xlabel('feature1')
plt.ylabel('feature2')
plt.show()

```

Points:

- For streaming/online needs, LOF is less suitable as-is (it's batch).
- For many features, reduce dimensionality for visualization (PCA).

Question 10 — Real-time energy demand forecasting & anomaly detection workflow

Paragraph:

For 15-minute streaming energy usage per region with weather and timestamps, you combine streaming preprocessing, online/rolling anomaly detection, short-term forecasting, and continuous monitoring. The solution uses sliding windows to re-train or update models, flags anomalies for operations, and feeds forecasts to scheduling/dispatch systems.

Short workflow (points):

1. Ingestion & preprocessing

- Collect 15-min records (timestamp, region, usage, weather, etc.) into a message queue (Kafka) or streaming pipeline.
- Clean, impute missing values, align timestamps, convert categorical features (region) to encodings, create features: lag values, rolling means, time features (hour, weekday), weather features, holiday flags.

2. Feature engineering (real-time)

- Maintain rolling features: 1-hr avg, 24-hr lag, same-time-last-week, rolling std.
- Normalize per region using rolling statistics to handle non-stationarity.

3. Anomaly detection in streaming

- **Technique choice:** use a hybrid approach: lightweight real-time detectors for edge/streaming (z-score on rolling windows, EWMA control charts) and periodic batch Isolation Forest or LOF on recent windowed data (e.g., last 7 days) for more robust detection.
- **Why:** Isolation Forest is robust and fast for batch windows; LOF picks local density anomalies; simple z/EWMA detects sudden spikes/drops immediately. Use ensemble voting to reduce false positives.

4. Short-term forecasting

- **Model choice:** SARIMAX (seasonal) or SARIMAX with exogenous regressors (weather, temperature) — i.e., **SARIMAX**. For ultra-low-latency forecasting, use light gradient-boosting regressors or online models with window retraining. SARIMAX is preferred when seasonality (daily/weekly) and exogenous variables matter.

5. Validation & monitoring

- Backtest with rolling-origin cross-validation (simulate live forecasts).
- Use metrics: MAE, RMSE for forecasts; precision/recall/F1 for anomaly alerts (if labeled events exist).
- Monitor model drift: track prediction error distributions and feature distribution drift (Population Stability Index).
- Alert / retrain policy: trigger retrain when error or drift exceeds threshold.

6. Operationalization

- Stream predictions/alerts to dashboards, auto-notify on critical anomalies, and integrate forecasts into dispatch and load-shedding decisions.

- Keep human-in-the-loop for triage during unusual events (storms, outages).

Example code (simplified streaming pseudo-implementation):

```
# Q10: Simplified streaming-pattern code (batch-window approach)
import pandas as pd
import numpy as np
from sklearn.ensemble import IsolationForest
from statsmodels.tsa.statespace.sarimax import SARIMAX

# pseudo function called every 15-minutes with a new batch/window
def process_window(window_df, model_state):
    """
    window_df: pandas DataFrame for the latest N*15min rows for one region
    model_state: dict with stored models, scalers, history
    """

    # 1. feature engineering
    df = window_df.copy()
    df['timestamp'] = pd.to_datetime(df['timestamp'])
    df.set_index('timestamp', inplace=True)
    df = df.asfreq('15T')
    df['hour'] = df.index.hour
    df['weekday'] = df.index.weekday
    # lags
    df['lag_1'] = df['usage'].shift(1)
    df['lag_96'] = df['usage'].shift(96) # same time previous day (96 *
15min = 24h)
    df['rolling_mean_4'] = df['usage'].rolling(4).mean() # last hour avg

    # 2. anomaly detection (quick z-score)
    recent = df['usage'].dropna()
    z = (recent - recent.rolling(96).mean()) / recent.rolling(96).std()
    immediate_anomalies = z[ (z.abs() > 4) ] # extreme spike/drop

    # 3. periodic batch IsolationForest on features (run every hour or n
windows)
    features = df[['usage', 'lag_1', 'lag_96', 'rolling_mean_4']].dropna()
    if len(features) > 200:
        iso = IsolationForest(contamination=0.01, random_state=0)
        iso.fit(features)
        preds = iso.predict(features)
        iso_anomalies = features[preds == -1]

    # 4. forecasting with SARIMAX for short-term (example: next 4 periods =
1 hour)
```

```

# keep a rolling history in model_state for this region
history = model_state.get('history', recent)
try:
    sar = SARIMAX(history, order=(1,0,1), seasonal_order=(1,1,1,96))
    res = sar.fit(dispatch=False)
    forecast = res.get_forecast(steps=4).predicted_mean
except Exception as e:
    forecast = pd.Series([history.iloc[-1]]*4) # fallback naive

# 5. summary output
return {
    'immediate_anomalies': immediate_anomalies,
    'iso_anomalies': iso_anomalies if 'iso_anomalies' in locals() else
pd.DataFrame(),
    'forecast_next_1h': forecast
}

```

Validation & Monitoring (summary):

- Use rolling backtesting to compute forecast error over many windows.
- Track false positive rate for anomaly system; tune thresholds.
- Deploy metrics to dashboards (Grafana) and set alerts on metric drift.
- Keep retrain cadence (daily/weekly) depending on drift.

How it helps business/operations:

- Early detection of abnormal spikes prevents outages and allows quicker dispatch of backup generation.
- Accurate short-term forecasts optimize dispatch planning, reduce operational costs, and improve grid stability.
- Automated alerts reduce manual monitoring load and enable targeted human intervention.